

Weka Sentiment Analysis Extension

Alexander Spivey

The Dataset: Wine Review dataset

```
df = pd.read_csv('winemag-data_first150k.csv')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150930 entries, 0 to 150929
```

```
Data columns (total 11 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	150930 non-null	int64
1	country	150925 non-null	object
2	description	150930 non-null	object
3	designation	105195 non-null	object
4	points	150930 non-null	int64
5	price	137235 non-null	float64
6	province	150925 non-null	object
7	region_1	125870 non-null	object
8	region_2	60953 non-null	object
9	variety	150930 non-null	object
10	winery	150930 non-null	object

```
dtypes: float64(1), int64(2), object(8)
memory usage: 12.7+ MB
```

	Unnamed: 0	country	description	designation	points	price	province	region_1	region_2	variety	winery
0	0	US	This tremendous 100% varietal wine hails from ...	Martha's Vineyard	96	235.0	California	Napa Valley	Napa	Cabernet Sauvignon	Heitz
1	1	Spain	Ripe aromas of fig, blackberry and cassis are ...	Carodorum Selección Especial Reserva	96	110.0	Northern Spain	Toro	NaN	Tinta de Toro	Bodega Carmen Rodríguez
2	2	US	Mac Watson honors the memory of a wine once ma...	Special Selected Late Harvest	96	90.0	California	Knights Valley	Sonoma	Sauvignon Blanc	Macauley
3	3	US	This spent 20 months in 30% new French oak, an...	Reserve	96	65.0	Oregon	Willamette Valley	Willamette Valley	Pinot Noir	Ponzi
4	4	France	This is the top wine from La Bégude, named aft...	La Brûlade	95	66.0	Provence	Bandol	NaN	Provence red blend	Domaine de la Bégude

```
df.dropna(inplace = True)
```

```
dropCols = ['Unnamed: 0', 'country', 'designation', 'price', 'province', 'region_1', 'region_2', 'variety', 'winery']
df.drop(dropCols, axis = 1, inplace = True)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 39241 entries, 0 to 150916
```

```
Data columns (total 2 columns):
```

#	Column	Non-Null Count	Dtype
0	description	39241 non-null	object
1	points	39241 non-null	int64

```
dtypes: int64(1), object(1)
memory usage: 919.7+ KB
```

Tokenizing the Entries

```
In [10]: import nltk
nltk.download('stopwords')

[nltk data] Downloading package stopwords to /home/aspiv/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
Out[10]: True
```

```
In [11]: from nltk.stem.porter import PorterStemmer
from nltk.corpus import stopwords

stop = stopwords.words('english')
porter = PorterStemmer()

def tokenize_stemmer(t):
    return [porter.stem(word) for word in t.split()]
```

```
In [12]: tokenize_stemmer('This is a test set of words that should catch stems')
```

```
Out[12]: ['thi',
'is',
'a',
'test',
'set',
'of',
'word',
'that',
'should',
'catch',
'stem']
```

```
In [13]: def tokenize_stemmer(t):
    l1 = [porter.stem(word) for word in t.split()]
    return [w for w in l1 if w not in stop]

tokenize_stemmer('This is a test set of words that should catch stems')
```

```
Out[13]: ['thi', 'test', 'set', 'word', 'catch', 'stem']
```

Result

```
In [19]: clf = LogisticRegressionCV(cv=5,  
                                   scoring = 'accuracy',  
                                   random_state = 101,  
                                   n_jobs=-1,  
                                   verbose=3,  
                                   max_iter=200).fit(X_train, y_train)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
```

```
RUNNING THE L-BFGS-B CODE
```

```
* * *
```

```
Machine precision = 2.220D-16
```

```
N =      592431      M =      10
```

```
At X0      0 variables are exactly at the bounds
```

```
At iterate   0   f=  5.73436D+04   |proj g|=  1.48010D+03
```

```
ITERATION    1
```

```
----- CAUCHY entered-----
```

```
There are      0   breakpoints
```

```
*** found in this segment
```

```
In [20]: # did only 200 iterations. when doing 300, i was sitting here for over half an hour then crash.
```

```
In [21]: clf.score(X_test, y_test)
```

```
Out[21]: 0.5313117156144486
```

```
In [24]: clf.score(X_train, y_train)
```

```
Out[24]: 0.9914627930682977
```

Double checking our Result - Random Decision Forest

```
In [18]: from sklearn.feature_extraction.text import TfidfVectorizer

cv = CountVectorizer()
rf = RandomForestClassifier(class_weight="balanced")
n_features = np.arange(10000, 30001, 10000)
def nfeature_accuracy_checker(vectorizer=cv, n_features=n_features, stop_words=None, ngram_range=(1, 1), classifier=rf):
    result = []
    print(classifier)
    print("\n")
    for n in n_features:
        vectorizer.set_params(stop_words=stop_words, max_features=n, ngram_range=ngram_range)
        checker_pipeline = Pipeline([
            ('vectorizer', vectorizer),
            ('classifier', classifier)
        ])
        print("Test result for {} features".format(n))
        nfeature_accuracy = accuracy_summary(checker_pipeline, X_train, y_train, X_test, y_test)
        result.append((n, nfeature_accuracy))
    return result
tfidf = TfidfVectorizer()
print("Result for trigram with stop words (Tfidf)\n")
feature_result_tgt = nfeature_accuracy_checker(vectorizer=tfidf, ngram_range=(1, 3))
```

Result for trigram with stop words (Tfidf)

RandomForestClassifier(class_weight='balanced')

Test result for 10000 features
accuracy score: 53.49%
Test result for 20000 features
accuracy score: 52.97%
Test result for 30000 features
accuracy score: 53.27%